# Agentic Appeals Flow — Phase 1 MVP

Azure AI Foundry · 3 Agents · Human-in-the-Loop · Cosmos DB Queue · AI Search RAG

**MANUAL PROVISIONING — FAST MVP FOR STAKEHOLDER DEMO**

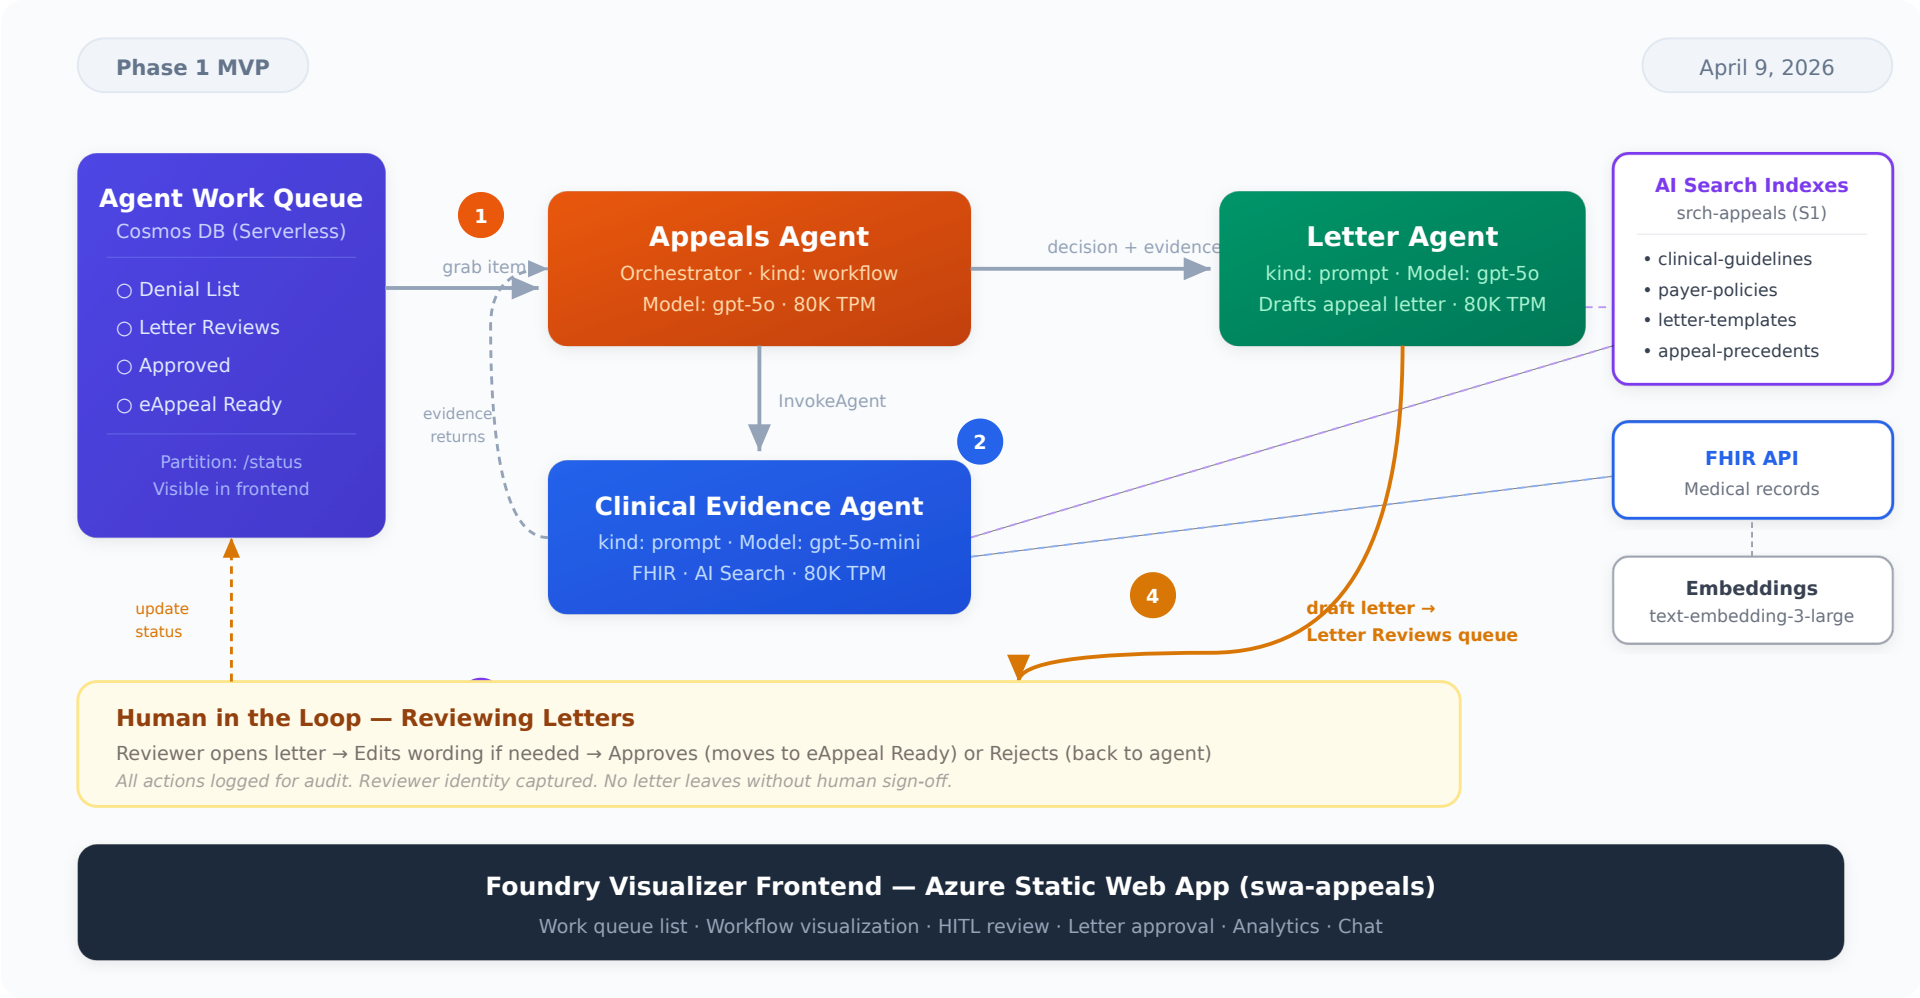


AI Foundry / Agents · Data & Storage · Monitoring · Frontend & Config · Human-in-the-Loop

# Azure Resources Required

Provision in order — later resources reference earlier ones. All resources in the existing MLAI Resource Group.

| # | RESOURCE | NAME | SKU / TIER | PURPOSE |
|---|---|---|---|---|
| **AI FOUNDRY — CORE AI PLATFORM** | | | | |
| 1 | **Azure AI Foundry Account** | `foundry-appeals` | S0 · AIServices | System-assigned managed identity. Custom domain. Houses all agents + model deployments. |
| 2 | **Foundry Project** | `appeals-project` | Workspace | Agents, connections, and deployments live here. Created inside the Foundry portal. |
| 3 | **Model: gpt-5o** | `gpt-5o` | Standard · 80K TPM | Appeals Agent (orchestrator) and Letter Agent. Strong reasoning + writing. |
| 4 | **Model: gpt-5o-mini** | `gpt-5o-mini` | Standard · 80K TPM | Clinical Evidence Agent. Cost-effective for retrieval-heavy work. |
| 5 | **Embeddings Model** | `text-embedding-3-large` | Standard · 350K TPM | Vector embeddings for AI Search RAG indexes. Runs during index build + every query. |
| **DATA & STORAGE** | | | | |
| 6 | **Azure AI Search** | `srch-appeals` | Standard S1 | 4 RAG indexes: clinical-guidelines, payer-policies, appeal-precedents, letter-templates. |
| 7 | **Cosmos DB** | `cosmos-appeals` | Serverless | Agent Work Queue + appeal records. Containers: work-queue, appeals, audit-log. |
| 8 | **Blob Storage** | `appeals-docs` | Hot tier | Raw denial documents, generated letters, uploaded evidence files. |
| **MONITORING** | | | | |
| 9 | **Application Insights** | `appi-appeals` | Pay-as-you-go | Agent trace events, Function logs, frontend telemetry. OpenTelemetry export from Foundry. |
| **FRONTEND & CONFIGURATION** | | | | |
| 10 | **Azure Static Web App** | `swa-appeals` | Free tier | React SPA + managed Functions proxy. 100 GB bandwidth, managed SSL, global CDN. |
| 11 | **Foundry Connections** | `ai-search-connection`<br>`cosmos-connection`<br>`blob-connection` | 3 connections | Agents reference connections by name — no raw endpoints or keys in agent config. |
| 12 | **SWA App Settings** | `FOUNDRY_ENDPOINT`<br>`FOUNDRY_PROJECT`<br>`FOUNDRY_API_KEY`<br>`APPINSIGHTS_CONN_STR` | 4 env vars | Runtime configuration for the managed Functions proxy. Set in SWA Configuration panel. |

**Foundry Connections:** After provisioning resources, create connections inside the Foundry project (AI Foundry portal → Project → Settings → Connections). Agents reference connections by name — they don't need raw endpoints or keys.

**Note:** Terraform and ADO pipelines come later. Manual portal provisioning gets the team unblocked fastest.